

ChatGPT Voice Discovery — Working Prototype

Implementation handoff for codebase / coding agent

1. Product intent

Build a focused Voice discovery enhancement inside a familiar ChatGPT mobile interaction model. This is not a redesign and not a separate Voice app.

Reference prototype: <https://chatgptvoiceusage.lovable.app>

The M3 PRD is the source of truth for product scope and behavior. The Lovable prototype is the visual and interaction reference.

Lovable prototype URL: <https://chatgptvoiceusage.lovable.app>

2. Critical requirement: real conversation

- The prototype must be genuinely interactive, not a scripted or prepopulated simulation.
- When a reviewer initiates Voice, the system must wait for the reviewer to actually speak.
- The reviewer's actual speech must be transcribed and displayed as the user message.
- The AI response must be generated dynamically from the actual user input and conversation history.
- The AI response should be spoken aloud.
- The reviewer must be able to continue the conversation naturally for multiple turns.
- Do not hard-code sample user messages, scripted assistant responses, or a fixed 3–4 turn conversation.
- The AI should respond to what the reviewer actually says. Do not force every response to end with a question.

3. Reviewer experience

ChatGPT home — familiar — Use a familiar ChatGPT-style mobile layout: top bar, empty state, standard composer. On desktop, center the experience in a phone-width column.

Voice discovery layer — Show the heading "What are you trying to figure out?" and three situation cards:

- Presentation in 30 minutes? — Practice your presentation through a natural conversation.
- Interview coming up? — Practice your answers through a natural conversation.
- Stuck between two options? — Talk it through and make your decision.

Hint: "Speak naturally. Use the language that feels comfortable."

Selecting a situation reveals the primary CTA "  Tap to Talk" and secondary option "Or type instead".

Voice conversation — inside the chat thread

The selected situation becomes the conversation context/title. The screen may show a short opening invitation, but this is NOT a simulated conversation turn.

- Presentation: "Tell me what you're presenting and what you want to practice."
- Interview: "Tell me about the role you're interviewing for."
- Decision: "Tell me what you're trying to decide between."

After that opening invitation, the conversation begins with the reviewer's real speech. Messages appear as normal ChatGPT-style bubbles. A single voice control communicates the current state: Ready to talk / Listening / Processing / Speaking, with an End action.

Feedback

- After End: "Did Voice help?" — Yes / Somewhat / No
- Then: "Would you use Voice for something similar again?" — Definitely / Probably / Probably not

4. Context and conversation behavior

The selected situation should provide context to the AI, but should not constrain the reviewer to a scripted flow.

- Presentation: help the reviewer think through presentation content, structure, delivery, questions, or concerns based on what they say.
- Interview: respond to the actual role, experience, question, concern, or answer the reviewer provides.
- Decision: help the reviewer explore the options, criteria, trade-offs, and uncertainty they actually describe.
- Maintain conversation history across turns so responses are contextual.
- Keep responses concise enough for a natural voice back-and-forth, while still being thoughtful and useful.

5. Multilingual behavior

Support natural English, Hindi, Hinglish, and mixed-language interaction where the underlying speech and language services support it. Do not add a language-selection screen. The reviewer should be able to speak in the language mix that feels natural.

6. Functional architecture

The implementation should use a minimal architecture suitable for a portfolio prototype:

- Browser microphone → speech-to-text
- Transcript + conversation history + selected situation → secure backend/server function
- Backend → LLM for a contextual response
- LLM response → text-to-speech
- Spoken response → reviewer; conversation continues

Technology choices are open to the coding agent. Do not make the implementation dependent on Lovable-specific services unless there is a clear reason. API keys/secrets must remain server-side.

7. State to support

- Selected situation
- Current conversation
- Conversation messages
- Voice active/inactive
- Voice state: Ready / Listening / Processing / Speaking
- Feedback response

8. Explicitly out of scope

- Replacing typing with Voice
- Redesigning ChatGPT
- A separate Voice app or dashboard
- User profiles or persona databases
- Language databases or language-selection UI
- Analytics dashboards
- Authentication, payments, or production infrastructure unless technically required for the prototype
- Production-grade Voice infrastructure
- A scripted/demo conversation

9. Acceptance criteria

- The reviewer immediately recognizes the experience as ChatGPT-like.
- The three contextual discovery situations are visible.
- The primary discovery CTA is exactly:  Tap to Talk.
- The reviewer can initiate Voice using their microphone.
- No fake or prepopulated conversation turns appear.
- The reviewer's actual speech becomes the user message.

- The AI responds dynamically to the actual content.
- The AI response is spoken aloud.
- The reviewer can continue for multiple turns.
- Conversation context is preserved across turns.
- The reviewer can end the conversation and provide the two-step feedback.
- Typing remains available through 'Or type instead'.

10. Source of truth

Use the M3 PRD for product requirements, scope, selected opportunity, JTBD, solution logic, metrics, and non-goals. Use the Lovable prototype URL for visual/interaction reference. Do not treat intermediate wireframes as a pixel-perfect UI specification; preserve the product intent and familiar ChatGPT interaction model.